

# ECINN: Efficient Counterfactuals from Invertible Neural Networks

BMVC 2021 Submission # 570

## Abstract

Counterfactual examples identify how inputs can be altered to change the predicted class of a classifier, thus opening up the black-box nature of, *e.g.*, deep neural networks. We propose a method, ECINN, that utilizes the generative capacities of invertible neural networks for image classification to generate counterfactual examples efficiently. In contrast to competing methods that sometimes need a thousand evaluations or more of the classifier, ECINN has a closed-form expression and generates a counterfactual in the time of only two evaluations. Arguably, the main challenge of generating counterfactual examples is to alter only input features that affect the predicted outcome, *i.e.*, class-dependent features. Our experiments demonstrate how ECINN alters class-dependent image regions to change the perceptual and predicted class, producing more realistically looking counterfactuals three orders of magnitude faster than competing methods.

## 1 Introduction

A great effort has been devoted to open up the black-box nature of deep neural networks for computer vision. Among others, heatmaps [3], class-maximizing samples [29], and contrastive examples [9] have been proposed; we focus on the latter. Contrastive examples are also known as counterfactual examples, even though models do not possess any causal structure as described in [26].<sup>1</sup> We adopt the setting from [13] and consider the generic question, “For situation  $X$ , why was the outcome  $Y$  and not  $Z$ ?”. We provide a counterfactual example to give an explanation of the form “Had  $X$  been  $\hat{X}$ , then the outcome would have been  $Z$ .”

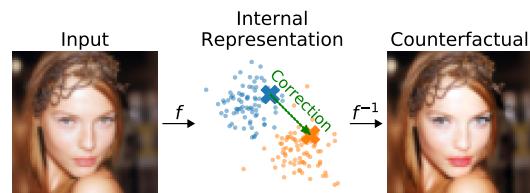


Figure 1:  $f$  transforms image *without* makeup (left) into internal representation which is corrected with closed-form expression (center).  $f^{-1}$  generates counterfactual example *with* makeup (right).

Being able to provide counterfactual examples for complex neural networks has an immense potential to improve human-model-interactions. To name but a few, surveillance systems could be assessed for biases when picking out candidates for screening and self-driving vehicles could be better diagnosed when misinterpreting their image feeds [13].

We propose Efficient Counterfactuals from Invertible Neural Networks (ECINN), which utilizes Invertible Neural Network (INN) classifiers to generate counterfactual examples.

<sup>043</sup> © 2021. The copyright of this document resides with its authors.

It may be distributed unchanged freely in print or electronic forms.

<sup>044</sup> <sup>1</sup>Counterfactual examples as described in [26] are based on structured causal graphs relating inputs and outputs.

<sup>045</sup> In the image domain, it is generally not known how to make such graphs and a causal analysis is thus not possible.

Figure 1 depicts the high-level structure of ECINN. An image of a woman *without* makeup (left) is transformed by an INN denoted  $f$  into an internal representation (center). The internal representation is corrected, as indicated by the green arrow, before being reverted by  $f^{-1}$  to form a counterfactual example *with* makeup (right).

The properties of INNs make a one-pass-solution possible. In contrast to usual discriminative models, INNs are known to have semantically organized latent spaces where translations in specific directions result in semantic changes in the input space [11]. Importantly, INNs even have full information-preservation between input and output layers in contrast to, e.g., auto-encoders [5], which allows exact recovery of inputs from outputs. As such, it can be argued that INNs are ideal for combining generative and discriminative capabilities [4].

Existing methods for generating counterfactual examples [1, 8, 9, 12, 13, 15, 25, 31, 32, 35] need to query the model under consideration many times due to various numerical optimization algorithms, obtain non-unique counterfactual examples, and need hyperparameter tuning. To the best of our knowledge, we introduce the first algorithm which uses only one forward and reverse pass, produces unique counterfactual examples, and needs no hyperparameter tuning. With just one forward and reverse pass, ECINN is even fast enough to be used in an interactive setting [7].

Good counterfactual examples are broadly agreed to be realistic, minimal, and actionable [12, 33]. In the image domain, however, minimal changes are hard to quantify in a semantically meaningful way. As such, we argue that the main challenge is to generate *realistically* looking images with *perceptible* changes only to class-relevant features.

We demonstrate experimentally how ECINN produces counterfactual examples leaving class-independent features largely untouched while class-dependent features are changed successfully. Experiments also demonstrate that our theoretically derived one-pass-solution yields running times more than three orders of magnitudes faster than competing methods.

## 2 Related Work

**Counterfactual Examples.** Many methods have been proposed for generating counterfactual examples or identifying counterfactual features. To name but a few, [1, 13, 31, 32, 34] operate on image data, [15, 16, 35] consider text, and yet other methods operate on relatively low dimensional data compared to images and text [8, 12, 33].

Methods for generating counterfactual examples can be categorized by the insights needed into the predictive model. Methods from the first category consider the predictive model as opaque and need no insight. Methods from the second category utilize gradients of the predictive model, while methods from the last category use internal data representations of the predictive model. All methods mentioned need to query the predictive model many times. In contrast, after a preprocessing step that needs to be done only once, our method uses a single forward and inverse pass through the model to generate a counterfactual example.

In the first category, methods operating on opaque models iteratively generate candidate sets before querying the predictive model to test candidates. [12] utilizes a greedy heuristic from simple data statistics to determine what input features to perturb, while [28] uses a genetic algorithm. [32] segments input images into super-pixels and use a greedy algorithm to perturb super-pixels. On text data, [35] finetunes a GPT-2 model [27] to generate similar sentences to the input sentence to generate new candidates. [34] identifies counterfactual regions in input images but does not generate counterfactual examples.

092 The second category employs gradient optimization techniques to generate counterfac-  
 093 tional examples. Albeit from a different perspective, previous work has developed methods for  
 094 synthesizing inputs that maximize desired (output) neurons. For example, [29] uses gradient  
 095 descent with an  $L_2$ -norm prior loss on a random input. [23] includes a local pixel varia-  
 096 tion prior to obtain more realistically looking features. [33] also proposed a different loss  
 097 based on the median absolute deviation. Even though the methods give insights into the inner  
 098 workings of the classifier, they suffer from generating unrealistic images. More recently, [25]  
 099 proposed to train a generative model to generate counterfactual examples. In a similar vein,  
 100 [9] utilizes a pretrained and fixed auto-encoder to identify latent codes that generate desired  
 101 outputs through gradient optimization. An extension of [9] is [31] which uses auto-encoders  
 102 or KD-trees to identify class prototypes which helps guide the gradient optimization. In  
 103 comparison to our one-pass-solution, the default maximum queries of the classifier in the  
 104 official code of [31] is 1000.<sup>2</sup> Finally, [22] uses gradients of the classifier to train an external  
 105 variational auto-encoder to generate counterfactuals fast. In contrast to ECINN, the method  
 106 has a substantial pre-computation time due to training of the auto-encoder.

107 The third category of methods contains two different strategies. First, [13] considers  
 108 convolutional neural networks as a composition of a (convolutional) feature extractor and a  
 109 classification network. They propose two algorithms for mixing feature fibers of the input  
 110 sample and a sample from the target class. Second, [1] similarly uses a part of the classifying  
 111 network as a feature extractor to cluster features, yielding an identification of semantic  
 112 features like stripes, wool, etc. A gradient descent algorithm successively adds or removes  
 113 features from the input to obtain a counterfactual example.

114 Although conceptually different, our work fits best into the third category. Instead of  
 115 generating counterfactual examples from an “arbitrary” neural network, we choose a specific  
 116 family of neural networks, INNs, to generate counterfactual examples efficiently without the  
 117 need for multiple queries of the model or memory consuming gradient computations.

118

119

120 **INNs as generative classifiers.** INNs have gained wide attention as unsupervised models  
 121 which allow generating realistically looking “fake” samples [10, 11, 19], typically referred  
 122 to as Normalizing Flows. Despite hidden in appendices, both [11] and [19] present sam-  
 123 ples generated from class-conditional INNs. Later, it was explicitly described how to com-  
 124 pose INNs with a Gaussian mixture model (GMM) to obtain a generative classifier [14, 24],  
 125 which both allows class-conditional sampling and classification. However, adding classifi-  
 126 cation abilities comes at a price. As demonstrated in [2], there is a trade-off between clas-  
 127 sification performance and the quality of the generated fake images. The work introduces  
 128 an information bottleneck loss, which explicitly trades off the classification and generation  
 129 performance through a hyperparameter  $\beta$ . [2] further introduces a new invertible model  
 130 architecture, which we refer to as IB-INN.

131 Regarding interpretability, [21] shows how conditional INNs can be trustworthy classi-  
 132 fiers by visualizing decision spaces, comparing class similarities, and computing posterior  
 133 heatmaps. In this work, we further show conditional INNs to be trustworthy by using them  
 134 for generating counterfactual examples.

135

---

136 <sup>2</sup>Official code: <https://docs.seldon.io/projects/alibi/en/stable/api/alibi-explainers.cproto.html>

### 3 Efficient Counterfactual Examples

This section constitutes our main contribution. We combine theoretical insights and practical observations from INNs to generate unique counterfactual examples from just one forward and inverse pass without the use of any numerical optimization techniques.

#### 3.1 Problem Statement

As mentioned, counterfactual examples indicate why an input was predicted to be one class rather than another. Specifically, we adopt the definition from [33] which states that counterfactual examples are statements taking the form: “Score  $p$  was returned because variables  $V$  had values  $(v_1, v_2, \dots)$  associated with them. If  $V$  instead had values  $(v'_1, v'_2, \dots)$ , and all other variables had remained constant, score  $p'$  would have been returned.” In the context of image classification, counterfactual examples are visualizations showing how the input image can be altered to change the predicted class.

**Desiderata.** In line with the desiderata of [12] and [33], we find that three properties are critical for counterfactuals to be useful. i) *Only semantically relevant features should be changed.* For example, facial features like lips and cheeks might change while the background should not when a counterfactual is generated for a face without makeup. ii) *Counterfactuals should look realistic.* Unrealistic counterfactuals might have misplaced eyes, extreme color values, or a “one-pixel-change” like the adversarial examples presented in [30]. iii) *Tipping-point counterfactuals and convincing counterfactuals should be prioritized.* With tipping-point, we refer to counterfactuals on the decision boundary, just where the prediction changes from the input to the target class and convincing counterfactuals are samples beyond the decision boundary that gets high probabilities for the target class.

**Definition 1.** (*tipping-point counterfactual*) Given a classifier with posterior probabilities  $p(y|x)$ , an input  $x \in \mathcal{X}$ , and a predicted class  $y = \arg \max_y p(y|x)$ , a counterfactual  $\hat{x}^{(q)}$  with target class  $q$  is a tipping-point counterfactual if there exists a path  $h : [0; 1] \rightarrow \mathcal{X}$  and constant  $C \in ]0; 1[$  such that  $h(0) = x$ ;  $h(C) = \hat{x}^{(q)}$ ; for  $c < C$ ,  $y$  has higher probability than  $q$ , i.e.,  $p(y|h(c)) > p(q|h(c))$ ; for  $c = C$ , probabilities are equal, i.e.,  $p(y|h(c)) = p(q|h(c))$ ; and for  $c > C$ ,  $q$  has higher probability than  $y$ , i.e.,  $p(y|h(c)) < p(q|h(c))$ .

**Definition 2.** (*convincing counterfactual*) given classifier  $p(y|x)$  for  $K$  classes and input  $x$  as defined above, a counterfactual  $\hat{x}^{(q)}$  is a convincing counterfactual if

$$\forall y' : y' \in \{1, \dots, K\} \setminus \{q\} \wedge p(q|\hat{x}^{(q)}) \gg p(y'|\hat{x}^{(q)}).$$

Tipping-point counterfactuals are essential because they represent minimal corrections to the input. However, they might not always make sense due to visual class differences. For example, when changing the predicted class of a cat to a dog, a tipping-point counterfactual might mix pointy and hanging ears because it is situated at the decision boundary. On the contrary, a convincing counterfactual would successfully show such transformation, but potentially with overly pronounced changes. Providing both types of explanations thus give a deeper insight into the decisions of the classifier.

In the supplementary material, we prove that ECINN produces valid tipping-point counterfactuals according to Definition 1 and in the experiments (Section 4), we verify that ECINN also complies with Definition 2 and the remaining desiderata.

### 3.2 Conditional INNs

We find INNs to be well suited for the counterfactual problem because they are bijective, *i.e.*, every latent vector corresponds to exactly one input. In contrast, typical classification models are inherently surjective, *i.e.*, there potentially exist many inputs which produce each output. In turn, INNs admits a single inverse pass to perfectly identify the right input while surjective models must rely on approximate solutions from less efficient numerical methods.

It is known that well-trained INNs have semantically organized latent spaces [11]. We believe that when many latent representations of samples from the same class are averaged, then class-independent information like background and object orientation will cancel out and leave just class-dependent information. ECINN isolates such latent class-dependent information to correct embeddings for generating counterfactual examples.

A conditional INN  $f$  is typically trained by computing latent vectors  $z = f(x)$  from input vectors  $x$  and using the latent vectors to fit a GMM to class labels  $y$ . However, to use  $z$  rather than  $x$  in the GMM, one must use the change-of-variables formula, which states that

$$\log p_X(x|y) = \log p_Z(f(x)|y) + \log |\det(J)|. \quad (1)$$

That is, the class-conditional log density of an input  $x$  in the image space,  $p_X(x|y)$ , is equal to the class-conditional log density of  $f(x)$  in the latent space  $p_Z(f(x)|y)$ , but with an additional Jacobian term,  $J = \frac{\partial f(x)}{\partial x}$ . We choose class-dependent latent densities to be Gaussians,  $p_Z(z|y) = \mathcal{N}(\mu_y, \mathbb{1})$ . By Bayes' rule, we notice that under a uniform prior distribution over labels,  $p(y) = \frac{1}{K}$  for  $K$  classes, the log posterior probability becomes

$$\log p_X(y|x) = \log \frac{p_X(x|y)}{\sum_{y'} p_X(x|y')} \propto -\|f(x) - \mu_y\|^2. \quad (2)$$

From Equation (2), we see that independent of the Jacobian determinant, latent vector  $z = f(x)$  will be predicted to be from the class  $y$  with the closest model mean,  $\mu_y$ . In turn, the latent space of the classifier can be analyzed under  $L_2$ -norms instead of less efficient and complex densities  $p_X(x|y)$ , which depend on the Jacobian determinant. In the following, we present how ECINN utilizes this insight to produce counterfactual examples efficiently.

### 3.3 ECINN

At a high level, ECINN transforms images into a latent space through an INN  $f$ . In the latent space, a closed-form expression changes the predicted class by correcting the embedding. From the corrected embedding, a counterfactual is generated by the inverse INN  $f^{-1}$ .

As a preprocessing step that needs to be done only once and takes just five seconds on MNIST, we group the training samples by their classified output,  $G_j = \{x | C(x) = j\}$ , where  $C(x) = \arg \max_y p_X(y|x)$  is the predicted class. Afterwards, we compute mean latent vectors  $\bar{\mu}_j = \frac{1}{|G_j|} \sum_{x \in G_j} f(x)$  for each class  $j$  and define the vector from  $\bar{\mu}_p$  to  $\bar{\mu}_q$  as  $\Delta_{p,q} = \bar{\mu}_q - \bar{\mu}_p$ .

Given a target class  $q$  and an input  $x$ , a counterfactual example  $\hat{x}^{(q)}$  is produced from the predicted class  $C(x) = p$  by adding a scaled version of  $\Delta_{p,q}$  to the latent space embedding  $z = f(x)$  and inverting it through the INN,

$$\hat{x}^{(q)} = f^{-1}(f(x) + \alpha \Delta_{p,q}). \quad (3)$$

It follows that a counterfactual example requires just one evaluation of  $f$  and  $f^{-1}$ .

To follow our third desideratum and provide both tipping-point and convincing counterfactuals, we compute two counterfactuals for each input with different values of  $\alpha$ . First, we choose  $\alpha_0$  to produce a tipping-point counterfactual. Due to Equation (2),  $\alpha_0$  is identified analytically such that  $\|z + \alpha_0 \Delta_{p,q} - \mu_p\| = \|z + \alpha_0 \Delta_{p,q} - \mu_q\|$ , which moves the latent vector to the decision boundary between the input and target class. The closed-form expression for  $\alpha_0$  is derived in the appendix (Section A), along with a proof that it complies with Definition 1 (Section B). Second, we choose  $\alpha_1$  such that the target class  $q$  is predicted with high confidence to produce a convincing counterfactual.  $\alpha_1$  is chosen heuristically to be  $\alpha_1 = \alpha_0 + \frac{4}{5}(1 - \alpha_0)$  (see supplementary material for details). Although not guaranteed that  $C(\hat{x}^{(q)}) = q$ , we observe that the relation holds in practice.

In Figure 2, we illustrate the intuition of ECINN. The figure shows two isotropic Gaussians in the latent space with a blue decision boundary. With green empty squares, we indicate the two computed means  $\bar{\mu}_p$  and  $\bar{\mu}_q$ , connected by  $\Delta_{p,q}$  (green arrow). The orange line passes through  $z$  in direction  $\Delta_{p,q}$ . The two points of interest are the blue square on the intersection of the blue and the orange line and the black square to the right. According to the model, the blue square constitutes a tipping-point counterfactual, and the black square is very likely to stem from class  $q$ , i.e., a convincing counterfactual.

In conclusion, we introduce ECINN which allows computing counterfactuals efficiently by utilizing theoretical and observational properties of INNs. ECINN complies with our first two desiderata by generating counterfactuals which represent class-dependents changes while leaving out most class-independent information. By providing both tipping-point and convincing counterfactuals, it also follows the third desideratum.

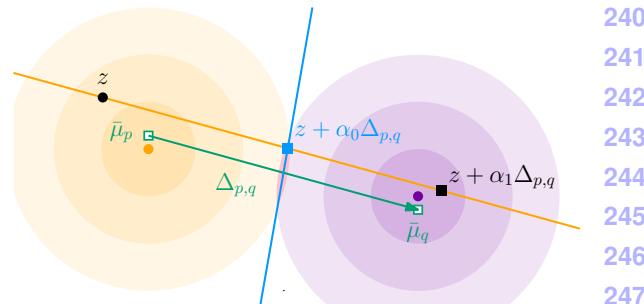


Figure 2: Latent space corrections by ECINN.

## 4 Experiments

In this section, we evaluate how our counterfactual examples perform. Our experiments show how ECINN produces meaningful counterfactual examples across three different image datasets, changes class-dependent features while maintaining class-independent features, and outperforms competing methods.

**Experimental Details.** We evaluate ECINN on a synthetic FakeMNIST dataset, the MNIST dataset [20], and the CelebA-HQ dataset [17]. On all three datasets, classification errors of the IB-INN models are comparable to those of a standard classification network (see Table 2 in the appendix). For all our experiments, we have trained IB-INN models “as-is.”<sup>3</sup> We found that  $\beta$ -values for IB-INN close to one strike a good balance between classification accuracy and generative performance (see appendix). We also provide an overview of hardware, all models used, hyperparameters, and the model performances in the appendix along with additional samples of all plots in the supplementary material. Results presented are all with samples from the test sets and were found to be consistent across samples.

<sup>3</sup>We adopted models and training code from <https://github.com/VLL-HD/IB-INN>.

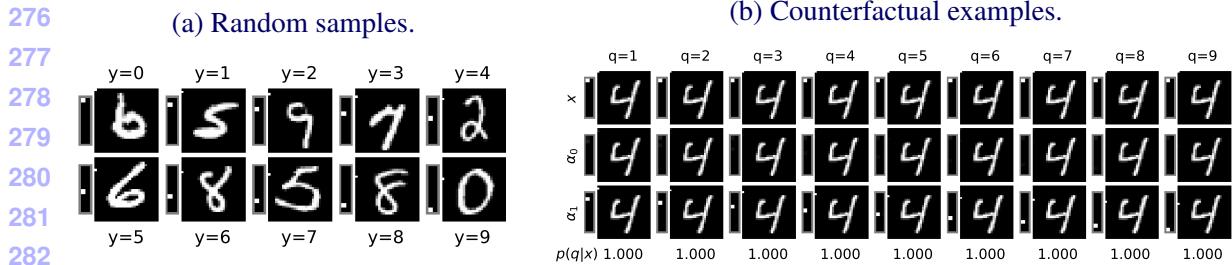


Figure 3: FakeMNIST dataset. For improved readability, smaller rectangles to the left of images magnify the top left  $10 \times 2$  pixels, indicating the class.

285

286 To further reproducibility, we provide anonymized code (`code.ipynb` in the supple-  
287 mentary material) which can be uploaded to Google Colab and run with one run command.  
288 Upon submission, we will release our code.

289

## 290 4.1 FakeMNIST 291

292 To verify ECINN in a controlled setting, we carefully design a dataset such that less than two  
293 percent of the pixels in each image are *class-dependent*. As argued, a proper counterfactual  
294 example for a well-trained model should alter only the class-dependent pixels and if no class-  
295 dependent information is present, each class should be equally likely.

296 The dataset is generated by randomly reassigning labels to images. We alter images *only*  
297 by injecting the information of the new labels in the top-left  $10 \times 1$  pixels; the  $i$ th top-left  
298 pixel will be white if the image is labeled  $i$ . For example, if an image gets label “5,” the  
299 sixth pixel in the left column is white. Figure 3a shows a sample from each of the ten classes.  
300 Only the top-left pixels depends on the labels; the depicted digits do not.

301 Figure 3b shows random sample from the class  $y = 0$  (first row) and tipping-point coun-  
302 terfactuals ( $\alpha_0$ ) in the second row. The third row includes convincing counterfactual exam-  
303 ples ( $\alpha_1$ ). Each column corresponds to a different target class  $q$ . Figure 3b shows that the  
304 dot in the top left corner of the input does change position, while the class-independent digit  
305 remains unchanged as expected. Specifically, the third row from left to right reveals how the  
306 dot in the top left corner travels downwards to end in the tenth pixel. The second row has no  
307 dot, which aligns well with the interpretation about equally likely class probabilities above.  
308

## 309 4.2 MNIST

310 Next, we apply ECINN to the MNIST dataset. First, we verify our second desideratum, *i.e.*,  
311 that ECINN produces realistic counterfactual examples. Second, we investigate how well  
312 class-independent features like font-weight and tilt are maintained by ECINN, *i.e.*, our first  
313 desideratum. Finally, we compare ECINN to two competing methods.  
314

315 **Realistic Counterfactuals.** In Figure 4,  
316 we depict counterfactual examples in the  
317 same fashion as Figure 3b. The figure  
318 shows how an image of a three is properly  
319 transformed into any of the remaining nine  
320 classes. Note that in the second row, the  
321 counterfactual examples are in many cases

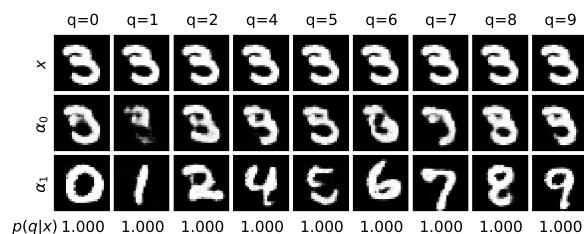


Figure 4: Same input, different targets.

such that even a human might mistake the image for both the input and target class. By contrast, the third row contains samples where the three has successfully transformed into the target class. This experiment demonstrates that ECINN complies with Definition 2 ( $p(q|x) = 1$  for all samples) and our second desideratum by generating realistic counterfactuals.

**Class-Independent Properties.** In Figure 5a and 5b, we demonstrate how class-independent properties like font-weight, tilt, and size are preserved during counterfactual generation. First, Figure 5a includes nine different inputs (first row), each from a different class, that are all translated to the target class,  $q = 0$ . We observe that the nine outcomes (row three) are perceptually different while resembling the target class. Each counterfactual example maintains class-independent properties from the input while resembling the target class. For example, the narrow and tilted one (first column) becomes a narrow and tilted zero. The observation suggests that ECINN maintains properties that are not directly dependent on the label.

In Figure 5b, we investigate how class-independent properties are maintained. We sample nine different images from the class  $y = 4$  and compute their counterfactual examples for the target class  $q = 9$ . We observe how bold inputs yield bold counterfactuals; likewise, slim inputs yield slim counterfactuals. Similar observations can be made for, e.g., tilt, size, and shape.

**Comparison.** In Figure 5c, we compare counterfactual examples generated with the algorithms proposed in [33] and [31] with our method.<sup>4</sup> Rows correspond to counterfactual methods and columns represents five different inputs. The figure shows that both competing methods generate more artificially looking counterfactual examples than ECINN. As the figure also shows, we found across many samples that counterfactuals generated by [33] and [31] generally look more artificial by having disconnected white pixels and being blurred, respectively. See supplementary material for additional samples.

In Table 1, we compare computation time on a single GPU, similar to [6]. For a fair comparison, we do not batch samples, as the framework for the competing methods does not support batching. The table shows how ECINN is more than  $6000\times$  faster than competing

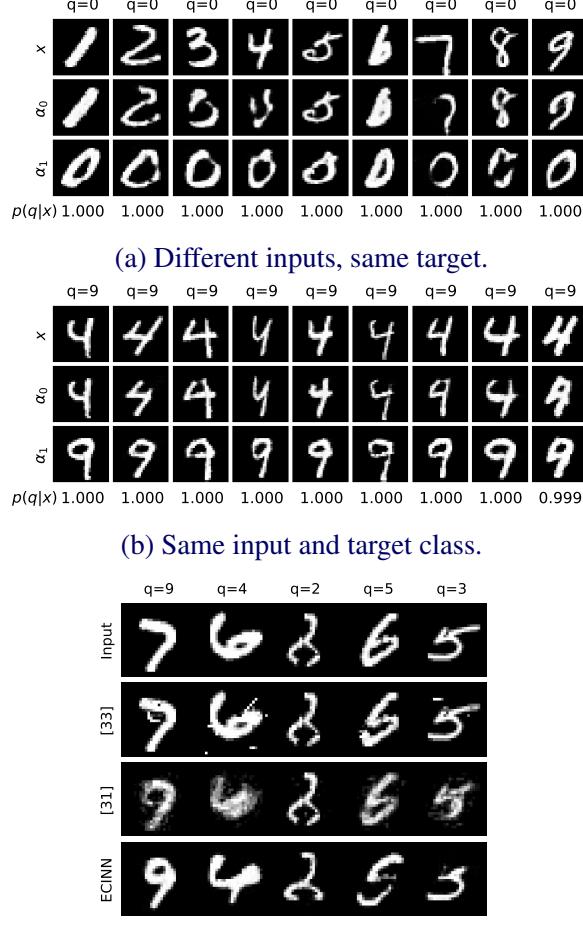


Figure 5: Counterfactuals for MNIST.

Mehod	Mean (std)	n
[33]	21.64 (7.99)	100
[31]	16.85 (0.35)	100
ECINN	<b>0.0025 (0.0002)</b>	$10^4$

Table 1: MNIST Computation times.

<sup>4</sup>Implementations found at <https://github.com/SeldonIO/alibi>; applied with default parameters.

322

323

324

325

326

327

328

329

330

331

332

333

334

335

336

337

338

339

340

341

342

343

344

345

346

347

348

349

350

351

352

353

354

355

356

357

358

359

360

361

362

363

364

365

366

367

368 methods. As it takes significantly less time than 0.1 second, ECINN can even be used in an  
 369 interactive setting [7], which is not possible with these competing methods.

370 In conclusion, we find that ECINN outperforms competing methods on both quality and  
 371 speed and comply with our desiderata by realistically changing the predicted and the per-  
 372 ceived class while maintaining class-independent features such as font-weight and tilt.

373

### 374 4.3 CelebA-HQ.

375

376 To evaluate ECINN on a more diverse and complex dataset, we extend our experiments to  
 377 the CelebA-HQ dataset. We train IB-INNs to predict various binary labels, where each class  
 378 is represented by at least 45% of the dataset.

379 In Figure 6, we show counterfactual examples for the smile versus frown label. The first  
 380 five columns depict how ECINN turns frowning people into smiling ones, while the last five  
 381 columns make smiling people frown. First, we observe that class irrelevant features such  
 382 as hair, skin color, and backgrounds remain perceptually unchanged as desired. Second,  
 383 we notice that some counterfactual examples in the last row look unrealistic. In particular,  
 384 it seems hard for the method to open and close mouths. In some cases, we also observe  
 385 small artifacts like the ones in the left-most pixels of the second column. Based on our  
 386 MNIST experiments, which did not suffer from computational limitations, we believe that  
 387 scaling from the roughly 40 million parameters used to around 200 million (as is common  
 388 with previous work [19]) can remove the artifacts and generate higher quality counterfactual  
 389 examples. Furthermore, the low-resolution version of CelebA-HQ that we use due to limited  
 390 resources is arguably harder to synthesize than higher resolutions. For further verification of  
 391 our findings, we include plots for models trained on other labels in Section E of the appendix.

391

392

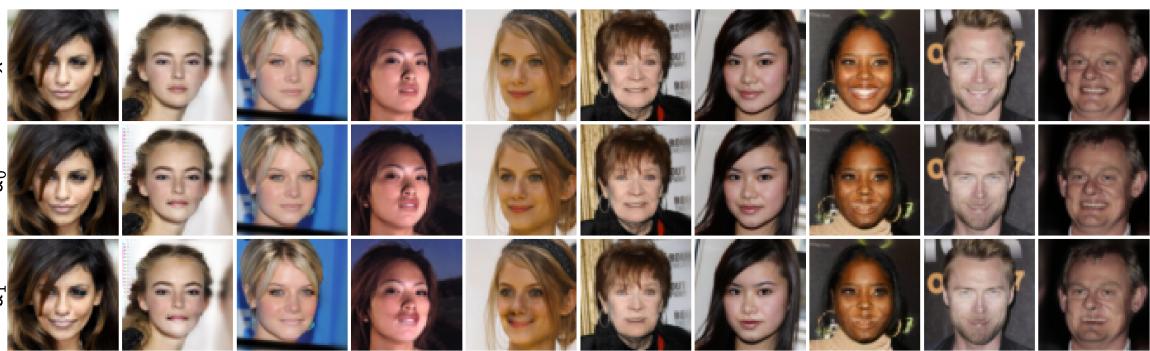
## 393 5 Conclusion

394

395 We introduce ECINN as an efficient method for computing counterfactual examples. Our  
 396 method is derived from theoretical and practical properties of a particular type of classifiers,  
 397 namely conditional INNs. While being three orders of magnitude faster than competing  
 398 methods, ECINN requires only one forward and one inverse pass, it generates a unique solu-  
 399 tion, and it requires no numerical optimization. In compliance with our desiderata, ECINN  
 400 generates counterfactual explanations that i) change only class-dependent features, ii) are  
 401 realistic, and iii) are diverse.

401

402



412 Figure 6: Counterfactual examples for frowning and smiling faces. First five columns have  
 413 target  $q = \text{smile}$  and last five columns  $q = \text{frown}$ .  $p(q|x) > 1 - 10^{-4}$  for all samples.

---

## References

- [1] Arjun Akula, Shuai Wang, and Song-Chun Zhu. CoCoX: Generating Conceptual and Counterfactual Explanations via Fault-Lines. *Proceedings of the AAAI Conference on Artificial Intelligence*, 2020.
- [2] Lynton Ardizzone, Radek Mackowiak, Carsten Rother, and Ullrich Köthe. Training normalizing flows with the information bottleneck for competitive generative classification. *NeurIPS*, 2020.
- [3] Sebastian Bach, Alexander Binder, Grégoire Montavon, Frederick Klauschen, Klaus Robert Müller, and Wojciech Samek. On pixel-wise explanations for non-linear classifier decisions by layer-wise relevance propagation. *PLoS ONE*, 2015.
- [4] Jens Behrmann, Will Grathwohl, Ricky T.Q. Chen, David Duvenaud, and Jörn Henrik Jacobsen. Invertible residual networks. In *ICML*, 2019.
- [5] Yoshua Bengio and Yann LeCun. Auto-Encoding Variational Bayes. In *ICLR*, 2014.
- [6] Francesco Bodria, Fosca Giannotti, Riccardo Guidotti, Francesca Naretto, Dino Pedreschi, and Salvatore Rinzivillo. Benchmarking and survey of explanation methods for black box models. *arXiv preprint arXiv:2102.13076*, 2021.
- [7] Stuart K Card, George G Robertson, and Jock D Mackinlay. The information visualizer, an information workspace. In *Proceedings of the SIGCHI Conference on Human factors in computing systems*, pages 181–186, 1991.
- [8] Furui Cheng, Yao Ming, and Huamin Qu. DECE: decision explorer with counterfactual explanations for machine learning models. *IEEE Transactions on Visualization and Computer Graphics*, 27(2):1438–1447, 2021.
- [9] Amit Dhurandhar, Pin Yu Chen, Ronny Luss, Chun Chen Tu, Paishun Ting, Karthikeyan Shanmugam, and Payel Das. Explanations based on the Missing: Towards Contrastive Explanations with Pertinent Negatives. In *NeurIPS*, 2018.
- [10] Laurent Dinh, David Krueger, and Yoshua Bengio. NICE: Non-Linear Independent Components Estimation. In *ICLR (Workshop)*, 2015.
- [11] Laurent Dinh, Jascha Sohl-Dickstein, and Samy Bengio. Density estimation using real NVP. In *ICLR*, 2019.
- [12] Oscar Gomez, Steffen Holter, Jun Yuan, and Enrico Bertini. ViCE: Visual Counterfactual Explanations for Machine Learning Models. *International Conference on Intelligent User Interfaces, Proceedings IUI*, pages 531–535, 2020.
- [13] Yash Goyal, Ziyan Wu, Jan Ernst, Dhruv Batra, Devi Parikh, and Stefan Lee. Counterfactual Visual Explanations. In *ICML*, 2019.
- [14] Pavel Izmailov, Polina Kirichenko, Marc Finzi, and Andrew Gordon Wilson. Semi-supervised learning with normalizing flows. In *ICML*, 2020.
- [15] Alon Jacovi, Swabha Swayamdipta, Shauli Ravfogel, Yanai Elazar, Yejin Choi, and Yoav Goldberg. Contrastive Explanations for Model Interpretability. *arXiv preprint arXiv:2103.01378*, 2021.

- 
- 460 [16] Sin-Han Kang, Honggyu Jung, Dong-Ok Won, and Seong-Whan Lee. Counterfac-  
461 tual explanation based on gradual construction for deep networks. *arXiv preprint*  
462 *arXiv:2008.01897*, 2020.
- 463 [17] Tero Karras, Timo Aila, Samuli Laine, and Jaakko Lehtinen. Progressive Growing of  
464 GANs for Improved Quality, Stability, and Variation. In *ICLR*, 2018.
- 465
- 466 [18] Diederik P Kingma and Jimmy Ba. Adam: A Method for Stochastic Optimization.  
467 *arXiv preprint arXiv:1412.6980*, 2014.
- 468
- 469 [19] Diederik P. Kingma and Prafulla Dhariwal. Glow: Generative Flow with Invertible 1x1  
470 Convolutions. In *NeurIPS*, 2018.
- 471 [20] Yann LeCun and Corinna Cortes. MNIST handwritten digit database. 2010. URL  
472 <http://yann.lecun.com/exdb/mnist/>.
- 473
- 474 [21] Radek Mackowiak, Lynton Ardizzone, Ullrich Köthe, and Carsten Rother. Generative  
475 classifiers as a basis for trustworthy computer vision. *arXiv preprint arXiv:2007.15036*,  
476 2020.
- 477 [22] Divyat Mahajan, Chenhao Tan, and Amit Sharma. Preserving causal constraints in  
478 counterfactual explanations for machine learning classifiers. In *CausalML: Machine*  
479 *Learning and Causal Inference for Improved Decision Making Workshop, NeurIPS*  
480 2019, December 2019.
- 481
- 482 [23] Anh Mai Nguyen, Jason Yosinski, and Jeff Clune. Deep neural networks are easily  
483 fooled: High confidence predictions for unrecognizable images. In *CVPR*, 2015.
- 484
- 485 [24] Tan M. Nguyen, Animesh Garg, Richard G. Baraniuk, and Anima Anandkumar. In-  
486 foCNF: Efficient conditional continuous normalizing flow using adaptive solvers. *arXiv*  
487 *preprint arXiv:1912.03978*, 2019.
- 488
- 489 [25] Jingjing Pan, Yash Goyal, and Stefan Lee. Question-conditioned counterfactual image  
490 generation for VQA. *arXiv preprint arXiv:1911.06352*, 2019.
- 491
- 492 [26] Judea Pearl. Causes of effects and effects of causes. *Sociological Methods & Research*,  
493 44(1):149–164, 2015.
- 494
- 495 [27] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya  
496 Sutskever. Language models are unsupervised multitask learners. *OpenAI blog*, 1  
497 (8):9, 2019.
- 498
- 499 [28] Shubham Sharma, Jette Henderson, and Joydeep Ghosh. CERTIFAI: Counterfactual  
500 Explanations for Robustness, Transparency, Interpretability, and Fairness of Artificial  
501 Intelligence models. *CoRR*, 2019. URL <http://arxiv.org/abs/1905.07857>.
- 502
- 503 [29] Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman. Deep inside convolutional  
504 networks: Visualising image classification models and saliency maps. In *ICLR*, 2014.
- 505
- 506 [30] Jiawei Su, Danilo Vasconcellos Vargas, and Kouichi Sakurai. One pixel attack for  
507 fooling deep neural networks. *IEEE Transactions on Evolutionary Computation*, 23  
508 (5):828–841, 2019.

- 
- [31] Arnaud Van Looveren and Janis Klaise. Interpretable counterfactual explanations guided by prototypes. *arXiv preprint arXiv:1907.02584*, 2019. 506  
507
- [32] Tom Vermeire and David Martens. Explainable image classification with evidence counterfactual. *arXiv preprint arXiv:2004.07511*, 2020. 508  
509  
510
- [33] Sandra Wachter, Brent Mittelstadt, and Chris Russell. Counterfactual explanations without opening the black box: Automated decisions and the GDPR. *Harv. JL & Tech.*, 31:841, 2017. 511  
512  
513
- [34] Pei Wang and Nuno Vasconcelos. SCOUT: Self-aware Discriminant Counterfactual Explanations. In *CVPR*, 2020. 514  
515  
516
- [35] Tongshuang Wu, Marco Tulio Ribeiro, Jeffrey Heer, and Daniel S. Weld. Polyjuice: Automated, General-purpose Counterfactual Generation. *arXiv preprint arXiv:2101.00288*, 2021. 517  
518  
519  
520  
521  
522  
523  
524  
525  
526  
527  
528  
529  
530  
531  
532  
533  
534  
535  
536  
537  
538  
539  
540  
541  
542  
543  
544  
545  
546  
547  
548  
549  
550  
551

552 **A Analytical  $\alpha$ -value,  $\alpha_0$**

553

554 Define  $y(\alpha) = z + \alpha\Delta_{p,q}$  to be the line intersecting  $z$  with direction  $\Delta_{p,q}$ . We wish to identify  
 555 the intersection between  $y(\alpha)$  and the hyperplane that constitutes the decision boundary  
 556 between the two normal distributions  $\mathcal{N}(\mu_p, \mathbb{1})$  and  $\mathcal{N}(\mu_q, \mathbb{1})$ . Due to the identity covariance  
 557 matrices of the Gaussians, we can define  $w = \mu_q - \mu_p$  and  $b = -\left(\frac{\mu_p + \mu_q}{2}\right)^T w$  to form the  
 558 decision boundary

559

$$w^T x + b = 0. \quad (4)$$

560

561 Equation (4) corresponds to the blue line in Figure 2.

562

563 To find the  $\alpha$ -value which corresponds to the intersection, set  $x = z + \alpha\Delta_{p,q}$  and solve  
 564 for  $\alpha$  in Equation (4):

565

$$w^T(z + \alpha\Delta_{p,q}) + b = 0 \quad (5)$$

566

$$\alpha w^T \Delta_{p,q} = -(w^T z + b) \quad (6)$$

567

$$\Rightarrow \alpha = -\frac{w^T z + b}{w^T \Delta_{p,q}}. \quad (7)$$

568

569

570 **Choice of  $\alpha_1$  value** As described, we found  $\alpha_1 = \alpha_0 + \frac{4}{5}(1 - \alpha_0)$  to be an appropriate value  
 571 for generating convincing counterfactuals across the three datasets covered in this work.  
 572 That said,  $\alpha_1 = 1$  would probably also have worked out fine. However, the goal was to stay  
 573 as close as possible to  $\alpha_0$  to change as little as possible, while still generating convincing  
 574 counterfactuals.

575

576 To give the reader an idea of the effect, we plot counterfactuals for five different inputs  
 577 for varying values of  $t$  in Figure 7. In the plot,  $\alpha_1$  is determined as a function of  $t$  and  $\alpha_0$ :

578

$$\alpha_1 = \alpha_0 + t(1 - \alpha_0). \quad (8)$$

579

580

581

582

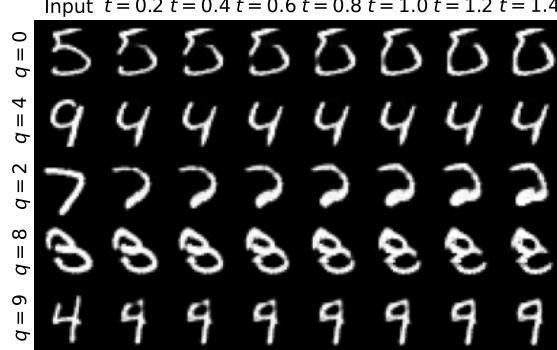
583

584

585

586

587



595

596

597 Figure 7: Effect of varying  $t$ , when generating counterfactuals using ECINN.

## B Tipping-point Counterfactuals

Here, we prove that ECINN produces tipping-point values according to Definition 1. Recall that we wish to find a constant  $C \in [0; 1]$  and  $h : [0; 1] \rightarrow \mathcal{X}$  such that  $p(y|h(c)) > p(q|h(c))$  whenever  $c < C$ , similarly  $p(y|h(c)) < p(q|h(c))$  whenever  $c > C$ , and finally  $p(y|h(c)) = p(q|h(c))$  when  $c = C$ . In the following, we choose  $C$  and  $h$  in such a way that they comply with Definition 1.

*Proof.* Let  $C = \frac{1}{2}$  and define  $\bar{c}$  as

$$\bar{c} = \begin{cases} 2\alpha_0 & \text{if } c \leq \frac{1}{2} \\ \alpha_0 + 2c - 1 & \text{otherwise} \end{cases}. \quad (9)$$

This way,  $\bar{c} < \alpha_0$  when  $c < C$  and  $\bar{c} > \alpha_0$  when  $c > C$ .

Now define  $h(c) = f^{-1}(z + \bar{c}\Delta_{y,q})$ , where  $f$  is the INN,  $z = f(x)$ , and  $\Delta_{y,q}$  is as defined in Section 3.3. Assume further that  $G(x) = y$  and  $G(f^{-1}(x + \Delta_{y,q})) = q$ , i.e., the input sample is correctly classified and the counterfactual is classified as class  $q$ .

*Sketch of proof:* we use the property of Equation (2) to show that  $\|\mu_y - z + c\Delta_{y,q}\| < \|\mu_q - z + c\Delta_{y,q}\|$  when  $c < C$  and vice versa.

By the change-of-variable formula (Equation (1)), Bayes' theorem, and the assumption  $p(y) = \frac{1}{K}$ , we have the relation

$$p(y|x) = \frac{p(f(x)|y)p(y)}{\sum_{y'} p(f(x)|y')p(y')} \quad (10)$$

$$= \frac{p(f(x)|y)}{\sum_{y'} p(f(x)|y')} \quad (11)$$

$$\Rightarrow \log p(y|x) = \log p_{z|y}(f(x)) + \log |det(J)| - \log \left[ \sum_{y'} p_{z|y}(f(x)|y') \right] - \log |det(J)| \quad (12)$$

$$= \log p_{z|y}(f(x)) - \log \left[ \sum_{y'} p_{z|y}(f(x)|y') \right] \quad (13)$$

and an identical relation holds for  $p(q|x)$

$$\log p(q|x) = \log p_{z|q}(f(x)) - \log \left[ \sum_{y'} p_{z|q}(f(x)) \right] \quad (14)$$

For a fixed  $x$ , we see that for  $\log p(y|x)$  to be greater than  $\log p(q|x)$ , only the first term matters. As  $p_{z|y} = \mathcal{N}(\mu_y, \mathbb{1})$ , we have that

$$\log p_{z|y}(z) = -\frac{d}{2} \log 2\pi - \frac{1}{2} \|\mu_y - z\|^2 \propto \|\mu_y - z\|^2, \quad (15)$$

and similarly for  $p_{z|q}$ . As such, by injecting  $h(c)$  into  $\log p_{z|y}$  it suffices to prove that  $\|\mu_y - z + \bar{c}\Delta_{y,q}\| = \|\mu_q - z + \bar{c}\Delta_{y,q}\|$  when  $c = C$ ,  $\|\mu_y - z + \bar{c}\Delta_{y,q}\| < \|\mu_q - z + \bar{c}\Delta_{y,q}\|$  when  $c < C$  and vice versa.

First, note that when  $c = C$ , then  $\bar{c} = \alpha_0$  so  $z + \bar{c}\Delta_{y,q} = z + \alpha_0\Delta_{y,q}$  and thus  $\|\mu_y - z + \bar{c}\Delta_{y,q}\| = \|\mu_q - z + \bar{c}\Delta_{y,q}\|$  holds by construction of  $\alpha_0$ .

644 Second, from the assumption that  $G(x) = y$ , we know that  $\|\mu_y - z + 0\Delta_{y,q}\| < \|\mu_q - z +$   
 645  $0\Delta_{y,q}\|$ . Similarly, from the assumption that  $G(f^{-1}(x + \Delta_{y,q})) = q$ , we know that  $\|\mu_y - z +$   
 646  $\Delta_{y,q}\| > \|\mu_q - z + \Delta_{y,q}\|$ . It follows, that when  $c < C$ ,  $\|\mu_y - z + c\Delta_{y,q}\| < \|\mu_q - z + c\Delta_{y,q}\|$  and  
 647 vice versa.  $\square$

648

## 649 C Experimental Details

651 In Table 2, we provide an overview of hyperparameters and performances of the networks  
 652 used in this work.

653

654 **IB-INN.** We have trained IB-INN models “as-is”<sup>5</sup> and adjusted only the  $\beta$ -value of the  
 655 loss function. On FakeMNIST and MNIST, the IB-INN models were trained for 60 epochs  
 656 with stochastic gradient descent and a milestone scheduler stepping from learning rate 0.07  
 657 to 0.007 after 50 epochs. On CelebA-HQ, the IB-INN models were trained for 800 epochs  
 658 with the Adam optimizer [18] and a milestone scheduler stepping with a factor  $\frac{1}{10}$  after every  
 659 200 epochs.

660

661

## 662 D IB-INN Model and Loss

663

664 The model architecture and loss function used in this work were proposed by [2]. The  
 665 loss was derived from an information bottleneck formulation with a hyperparameter,  $\beta$ , that  
 666 allows trading off generative and classification capabilities. The loss function is based on  
 667 mutual information  $I$ :

$$668 \quad \mathcal{L}_{IB} = I(X, Z) - \beta I(Z, Y). \quad (16)$$

669 Mutual information quantifies the amount of information which is shared between variables.<sup>6</sup>  
 670 As such, by minimizing  $\mathcal{L}_{IB}$ , the mutual information between the input and the latent vec-  
 671 tor is minimized while the mutual information between the latent vector and class label is  
 672 maximized. In practice, the first term,  $I(X, Z)$ , can be thought of as a generative loss, which  
 673 results in a good performance on generating images. The second term,  $I(Z, Y)$ , is closely

674

675 <sup>5</sup>IB-INN code: <https://github.com/VLL-HD/IB-INN>

676 <sup>6</sup>For an invertible mapping  $f$  and  $Z = f(X)$ ,  $\mathcal{L}_{IB}$  is, in fact, ill-defined, and the authors [2] add noise to  $X$  to  
 677 overcome the issue.

678

679	Dataset	$\beta$	BPD	Err.
680	FakeMNIST	1.4265	1.77	0%
681	MNIST	1.4265	1.89	0.85%
682 CelebA-HQ				
683	Smile	1	3.32	7.42%
684	High cheekbones	1	3.09	14.38%
685	Lipstick	1	3.06	4.87%
686	Heavy makeup	1	3.08	12.68%

688 Table 2: Hyperparameters, negative log-likelihood measured in bits per dimension (BPD),  
 689 and error rates for the models used in this work.

---

**16****AUTHOR(S): ECINN: EFFICIENT COUNTERFACTUALS FROM INNS**

---

related to the categorical cross-entropy loss, thus promoting high accuracy. Throughout our experiments, we use models trained with the IB-INN loss,  $\mathcal{L}_{IB}$ . 690  
691

For simplicity, we do not include experiments across multiple values of  $\beta$  in the main paper. Overall, we find that values close to one strike a good balance between counterfactual examples and model accuracy in our experiments. We do, however, include Figure 8 which demonstrates the conflicting effect of  $\beta$  on the quality of counterfactuals and the accuracy of the model. 692  
693  
694  
695  
696

697

698

699

700

701

702

703

704

705

706

707

708

709

710

711

712

713

714

715

716

717

718

719

720

721

722

723

724

725

726

727

728

729

730

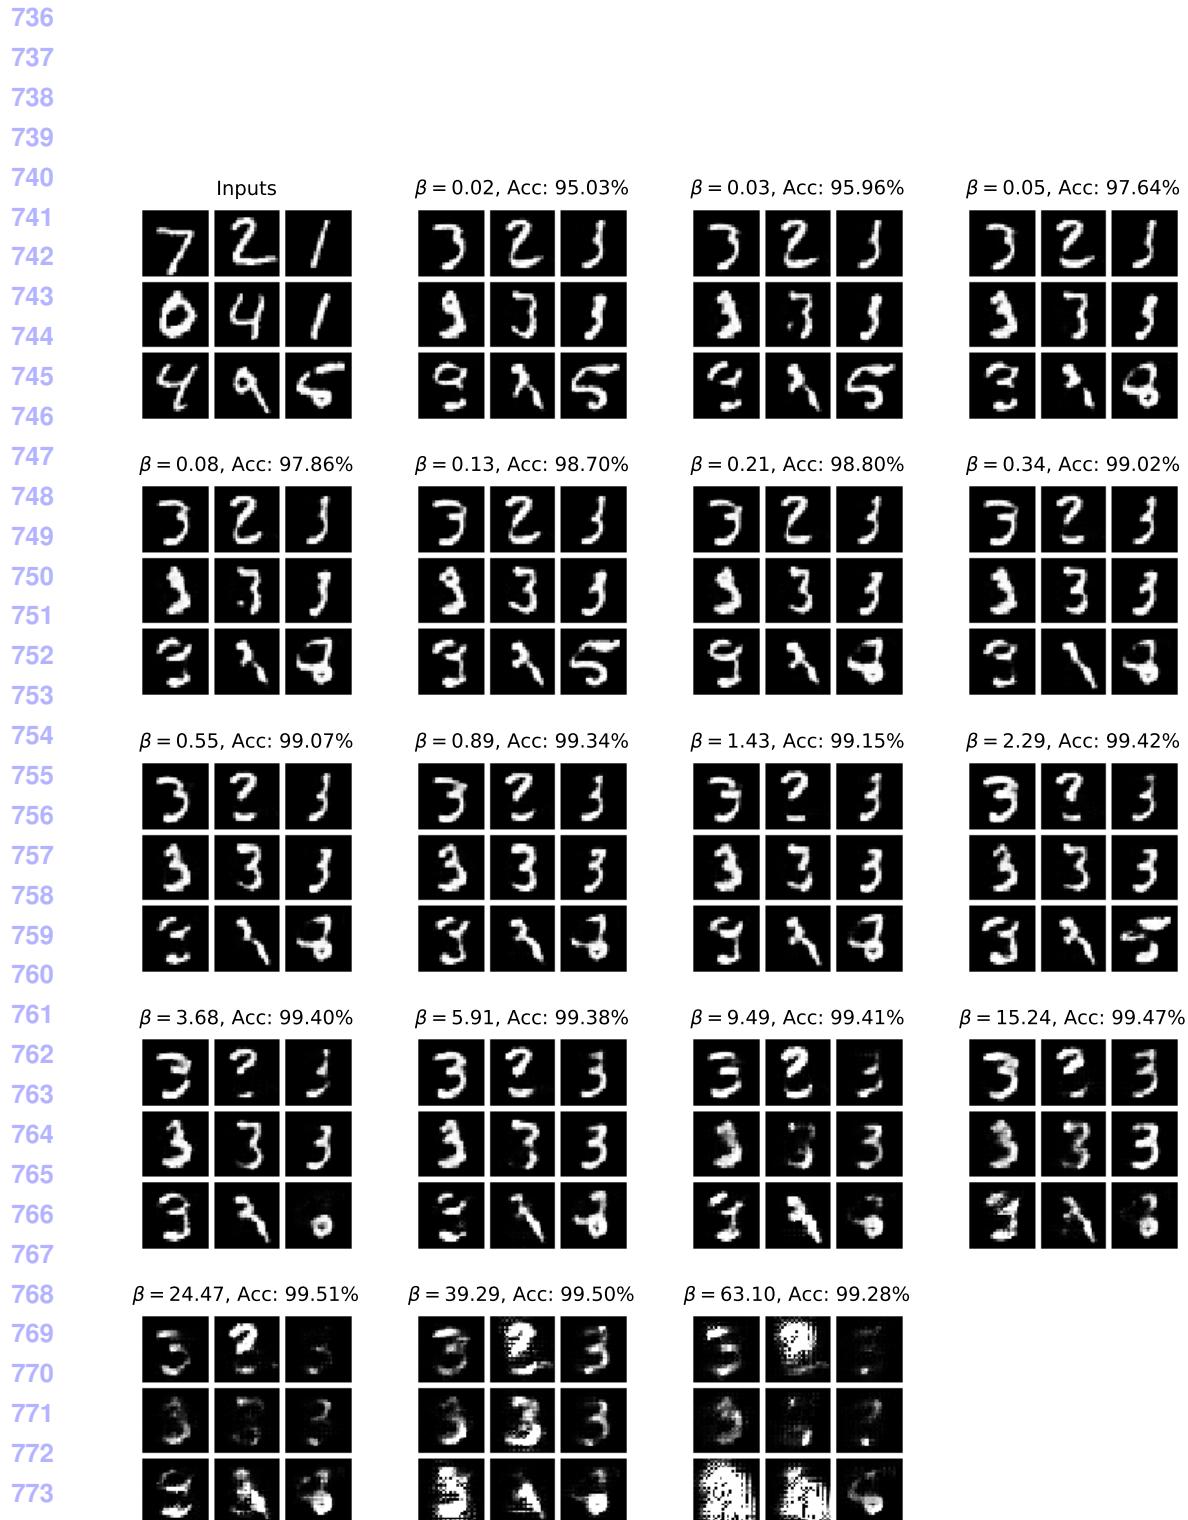
731

732

733

734

735

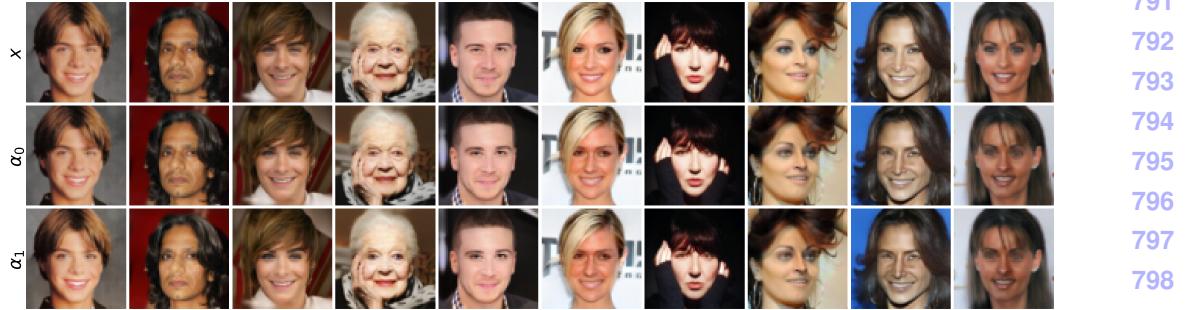


774 Figure 8: Counterfactual examples for MNIST models trained with different values of  $\beta$ .  
 775 The top left square represents the input images that are all changed with target  $q = 3$ . Above  
 776 plots are  $\beta$ -values in ascending order and corresponding test set accuracies.

777  
778  
779  
780  
781



(a) Label: High cheekbones.



(b) Label: Wearing lipstick.



(c) Label: Heavy makeup.

Figure 9: CelebA-HQ counterfactual examples. First five columns are inputs with negative labels and counterfactuals with positive labels and vice versa for the last five columns.

## E Additional Samples

In Figure 9, we include counterfactual examples similar to Figure 6 for three additional labels. We also include pdfs with extra samples of all figures from our experiments. For each figure, there is a corresponding pdf in the related work zip-file. For example, Figure 9a has a corresponding pdf in the supplementary material named `figure9a.pdf` with additional samples.

828 F Hardware Specifications

All experiments were run on a single machine learning server with 128GB system memory, an Intel(R) Xeon(R) Silver 4214 CPU @ 2.20GHz processor, and 5 NVIDIA RTX 2080 Ti. See all details below.

```
833 $ nvidia-smi -L
834 GPU 0: GeForce RTX 2080 Ti (UUID: ...)
835 GPU 1: GeForce RTX 2080 Ti (UUID: ...)
836 GPU 2: GeForce RTX 2080 Ti (UUID: ...)
837 GPU 3: GeForce RTX 2080 Ti (UUID: ...)
838
839 $ lscpu
840 Architecture:          x86_64
841 CPU op-mode(s):        32-bit, 64-bit
842 Byte Order:            Little Endian
843 CPU(s):                48
844 On-line CPU(s) list:  0-47
845 Thread(s) per core:   2
846 Core(s) per socket:   12
847 Socket(s):            2
848 NUMA node(s):          2
849 Vendor ID:             GenuineIntel
850 CPU family:            6
851 Model:                 85
852 Model name:           Intel(R) Xeon(R) Silver 4214 CPU @ 2.20GHz
853 Stepping:               7
854 CPU MHz:                1000.777
855 CPU max MHz:           2201.0000
856 CPU min MHz:           1000.0000
857 BogoMIPS:                4400.00
858 CPU Virtualization:    VT-x
859 CPU L1d cache:          32K
860 CPU L1i cache:          32K
861 CPU L2 cache:            1024K
862 CPU L3 cache:            16896K
863 NUMA node0 CPU(s):      0-11,24-35
864 NUMA node1 CPU(s):      12-23,36-47
865
866 $ lsblk
867
868
869
870
871
872
873
```